



[About Keras](#)

[Getting started](#)

[Developer guides](#)

Keras API reference

Models API

Layers API

Callbacks API

Data preprocessing

Optimizers

Metrics

Losses

Built-in small datasets

Keras Applications

Utilities

[Code examples](#)

[Why choose Keras?](#)

[Community & governance](#)

[Contributing to Keras](#)

» [Keras API reference](#) / Keras Applications

# Keras Applications

Keras Applications are deep learning models that are made available alongside pre-trained weights. These models can be used for prediction, feature extraction, and fine-tuning.

Weights are downloaded automatically when instantiating a model. They are stored at `~/.keras/models/`.

Upon instantiation, the models will be built according to the image data format set in your Keras configuration file at `~/.keras/keras.json`. For instance, if you have set `image_data_format=channels_last`, then any model loaded from this repository will get built according to the TensorFlow data format convention, "Height-Width-Depth".

## Available models

Model	Size	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth
<a href="#">Xception</a>	88 MB	0.790	0.945	22,910,480	126
<a href="#">VGG16</a>	528 MB	0.713	0.901	138,357,544	23
<a href="#">VGG19</a>	549 MB	0.713	0.900	143,667,240	26
<a href="#">ResNet50</a>	98 MB	0.749	0.921	25,636,712	-
<a href="#">ResNet101</a>	171 MB	0.764	0.928	44,707,176	-
<a href="#">ResNet152</a>	232 MB	0.766	0.931	60,419,944	-
<a href="#">ResNet50V2</a>	98 MB	0.760	0.930	25,613,800	-
<a href="#">ResNet101V2</a>	171 MB	0.772	0.938	44,675,560	-
<a href="#">ResNet152V2</a>	232 MB	0.780	0.942	60,380,648	-
<a href="#">InceptionV3</a>	92 MB	0.779	0.937	23,851,784	159
<a href="#">InceptionResNetV2</a>	215 MB	0.803	0.953	55,873,736	572
<a href="#">MobileNet</a>	16 MB	0.704	0.895	4,253,864	88
<a href="#">MobileNetV2</a>	14 MB	0.713	0.901	3,538,984	88
<a href="#">DenseNet121</a>	33 MB	0.750	0.923	8,062,504	121
<a href="#">DenseNet169</a>	57 MB	0.762	0.932	14,307,880	169
<a href="#">DenseNet201</a>	80 MB	0.773	0.936	20,242,984	201
<a href="#">NASNetMobile</a>	23 MB	0.744	0.919	5,326,716	-
<a href="#">NASNetLarge</a>	343 MB	0.825	0.960	88,949,818	-
<a href="#">EfficientNetB0</a>	29 MB	-	-	5,330,571	-
<a href="#">EfficientNetB1</a>	31 MB	-	-	7,856,239	-
<a href="#">EfficientNetB2</a>	36 MB	-	-	9,177,569	-
<a href="#">EfficientNetB3</a>	48 MB	-	-	12,320,535	-
<a href="#">EfficientNetB4</a>	75 MB	-	-	19,466,823	-
<a href="#">EfficientNetB5</a>	118 MB	-	-	30,562,527	-
<a href="#">EfficientNetB6</a>	166 MB	-	-	43,265,143	-
<a href="#">EfficientNetB7</a>	256 MB	-	-	66,658,687	-

The top-1 and top-5 accuracy refers to the model's performance on the ImageNet validation dataset.

Depth refers to the topological depth of the network. This includes activation layers, batch normalization layers etc.

# Usage examples for image classification models

## Classify ImageNet classes with ResNet50

```
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
import numpy as np

model = ResNet50(weights='imagenet')

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

preds = model.predict(x)
# decode the results into a list of tuples (class, description, probability)
# (one such list for each sample in the batch)
print('Predicted:', decode_predictions(preds, top=3)[0])
# Predicted: [(u'n02504013', u'Indian_elephant', 0.82658225), (u'n01871265', u'tusker',
0.1122357), (u'n02504458', u'African_elephant', 0.061040461)]
```

## Extract features with VGG16

```
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.vgg16 import preprocess_input
import numpy as np

model = VGG16(weights='imagenet', include_top=False)

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

features = model.predict(x)
```

## Extract features from an arbitrary intermediate layer with VGG19

```
from tensorflow.keras.applications.vgg19 import VGG19
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.vgg19 import preprocess_input
from tensorflow.keras.models import Model
import numpy as np

base_model = VGG19(weights='imagenet')
model = Model(inputs=base_model.input, outputs=base_model.get_layer('block4_pool').output)

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

block4_pool_features = model.predict(x)
```

## Fine-tune InceptionV3 on a new set of classes

```

from tensorflow.keras.applications.inception_v3 import InceptionV3
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D

# create the base pre-trained model
base_model = InceptionV3(weights='imagenet', include_top=False)

# add a global spatial average pooling layer
x = base_model.output
x = GlobalAveragePooling2D()(x)
# let's add a fully-connected layer
x = Dense(1024, activation='relu')(x)
# and a logistic layer -- let's say we have 200 classes
predictions = Dense(200, activation='softmax')(x)

# this is the model we will train
model = Model(inputs=base_model.input, outputs=predictions)

# first: train only the top layers (which were randomly initialized)
# i.e. freeze all convolutional InceptionV3 layers
for layer in base_model.layers:
    layer.trainable = False

# compile the model (should be done *after* setting layers to non-trainable)
model.compile(optimizer='rmsprop', loss='categorical_crossentropy')

# train the model on the new data for a few epochs
model.fit(...)

# at this point, the top layers are well trained and we can start fine-tuning
# convolutional layers from inception V3. We will freeze the bottom N layers
# and train the remaining top layers.

# let's visualize layer names and layer indices to see how many layers
# we should freeze:
for i, layer in enumerate(base_model.layers):
    print(i, layer.name)

# we chose to train the top 2 inception blocks, i.e. we will freeze
# the first 249 layers and unfreeze the rest:
for layer in model.layers[:249]:
    layer.trainable = False
for layer in model.layers[249:]:
    layer.trainable = True

# we need to recompile the model for these modifications to take effect
# we use SGD with a low learning rate
from tensorflow.keras.optimizers import SGD
model.compile(optimizer=SGD(lr=0.0001, momentum=0.9), loss='categorical_crossentropy')

# we train our model again (this time fine-tuning the top 2 inception blocks
# alongside the top Dense layers
model.fit(...)

```

## Build InceptionV3 over a custom input tensor

```

from tensorflow.keras.applications.inception_v3 import InceptionV3
from tensorflow.keras.layers import Input

# this could also be the output a different Keras model or layer
input_tensor = Input(shape=(224, 224, 3))

model = InceptionV3(input_tensor=input_tensor, weights='imagenet', include_top=True)

```